

Gold Standard Example 1

Example 1: Enterprise SaaS — Device Provisioning Wizard (8 screens, Tool-agnostic, Web)

PRODUCT BRIEF

HP WXP – enterprise device management platform for IT teams.

Platform: Web.

User: Adam Chen (IT Administrator) – automation-first, but needs evidence before trusting AI recommendations.

Style: Clean enterprise SaaS, similar to Datadog or HashiCorp.

Functional, data-dense, professional. Design system is pre-loaded.

Navigation: Left sidebar with Dashboard, Fleet, Policies, Patches, Security, Assets, Reports, Settings. Top bar with search, notifications bell, and user avatar. Fleet is active for this flow.

Flow: Zero-touch provisioning – 8 screens. Import devices, get AI profile recommendations, review and assign, then monitor deployment.

SCREEN 1 of 8: Fleet Dashboard

Adam opens the Fleet Dashboard, his home base for managing all company devices. He's here to start a new device provisioning batch.

The sidebar shows Fleet as the active section. The dashboard displays a priorities banner at the top with three actionable items:

- "75 devices pending deployment" with a link to start provisioning
- "12 devices offline >48h" flagged as critical
- "4 compliance failures" needing review

Below that, a row of 6 key metrics: 1,847 total enrolled devices (+24 this week), 7 active alerts (2 critical), 94.2% compliance (target 98%), DEX score 76 (healthy), 88% patch level, and 12 offline devices.

The lower section splits into an alerts table on the left showing the 5 most recent alerts with device ID, issue description, severity

(Critical, High, Warning), and timestamp. On the right, a system activity feed showing recent events including "Deployment created DEPL-2024-0847."

Adam clicks "+ Add Devices" to begin the provisioning wizard.

SCREEN 2 of 8: Import Device List

Adam has clicked "+ Add Devices" from the dashboard and is now on step 1 of the provisioning wizard. A 4-step progress indicator at the top shows: Import → AI Profiles → Assign → Review. Step 1 is active.

The screen offers two ways to import devices. On the left, a file upload area where Adam can drag a CSV file or browse to select one, with a link to download a CSV template. Below that, an alternative: enter a purchase order number. Adam has typed "PO-2024-11892" and the system confirms: "PO-2024-11892 matched · HP Direct Fleet Store · 75 × HP EliteBook 840 G11."

On the right, the import preview shows results: 75 devices imported, 71 valid, 4 with errors. A table lists device serial numbers and their validation status. Sample rows:

- 5CG4ABCD01 – Valid
- 5CG4ABCD02 – Valid
- 5CG4ABCD03 – Valid
- 5CG4QRST55 – Error: Unrecognized serial
- 5CG4MNOP33 – Error: Model mismatch
- 5CG4ABCD06 – Valid

A summary bar reads "71 valid · 4 errors · 0 duplicates."

The "Next: Get AI Profile Recommendations" button is at the bottom. It's disabled with tooltip "Import at least one device" when no valid devices exist. A Cancel button returns to the dashboard.

SCREEN 3 of 8: AI Profile Recommendations

Adam advances to step 2 of the wizard. The progress indicator shows step 1 (Import) complete with a checkmark, step 2 (AI Profiles) active.

The system has analyzed the 71 valid devices and identified 5 department groups. The left panel shows each group with a role dropdown:

- Finance: 22 devices, role set to "Finance / Legal"
- Strategy: 18 devices, role set to "Executive"
- Engineering: 14 devices, role set to "Developer / Engineering"
- Remote: 11 devices, role set to "Remote Worker"
- TBD: 10 devices – flagged with "ACTION REQUIRED" – the role dropdown is empty. "No HR match found in database."

The right panel shows AI-generated profile recommendations for each resolved group. Finance is expanded showing FIN-SEC-STRICT at 94% confidence, with three reasons: enhanced isolation for financial data, zero-trust network access by default, and real-time encryption monitoring. The other groups show collapsed rows with their matched profiles: Strategy → EXEC-PREMIUM (87%), Engineering → DEV-ELEVATED (91%), Remote → REMOTE-WORKER (89%).

An info banner at the bottom of the recommendations reads: "TBD group (10 devices) unresolved. AI cannot recommend profiles until a role is assigned."

The "Next" button is disabled with tooltip "Resolve TBD group first." Back returns to Import.

VARIATION – Build trust in AI recommendations

Screen 3: AI Profile Recommendations

This variation addresses: HMW help Adam trust AI-generated profile recommendations when he's naturally skeptical of automation?

Instead of showing recommendations as a simple list with confidence percentages, this version adds a prominent "How we chose this" section for each recommendation. When Adam expands any group, he sees: (1) the recommended profile, (2) a confidence score with a visual bar, (3) a "Reasoning" section showing the 3-4 data signals the AI used (e.g., "22 devices matched Finance department in HR database", "Historical provisioning for Finance used FIN-SEC-STRICT in 89% of cases"), and (4) a "Next best alternative" showing the

runner-up profile with a "Compare" button that opens a side-by-side comparison table.

Adam can also see a small "Override" button next to each recommendation, with reassurance text: "You can always change this in the next step."

All other content remains the same as the base prompt.

SCREEN 4 of 8: Review Profile Recommendations

Adam advances to step 3. Steps 1-2 show checkmarks, step 3 (Assign) is active.

This is a review gate – Adam must accept or override each AI recommendation before proceeding. A summary bar at the top shows "0 of 5 segments reviewed" with a progress indicator and an "Accept all recommendations" shortcut button.

The screen shows an accordion-style list of all 5 groups. Each row displays the group name, device count, recommended profile, match percentage, and a status (Pending/Accepted/Overridden).

One group (Engineering, 14 devices) is expanded, showing a comparison between the AI recommendation (DEV-ELEVATED – "Optimized for high-compute dev tasks" – with software chips: BitLocker, VS Code Insiders, Docker Desktop, Postman Enterprise, Git LFS) and the next-best alternative (GENERAL-STD – "Standard corporate profile for general productivity"). Adam can click "Accept recommendation" or "Override."

A footer note reads: "All decisions are logged to the audit trail."

The "Next: Assign Users" button is disabled until all 5 segments have been reviewed.

VARIATION – Make review faster without losing control
Screen 4: Review Profile Recommendations

This variation addresses: HMW ensure recommendations are transparent and adjustable without slowing Adam down?

Instead of requiring Adam to expand and explicitly accept each group, this version shows all 5 groups pre-accepted with a visual indicator (green checkmark), but with a prominent "Review changes" summary showing what was auto-accepted. Each group has an "I disagree" button that expands the comparison view. This inverts the interaction model: trust by default, override on exception – which matches Adam's "automation-first" behavioral trait while still giving him full control.

The summary bar reads "5 of 5 auto-accepted – review any you disagree with."

SCREEN 5 of 8: Assign Users to Devices

Adam is on step 3. He needs to match users to specific devices. The system has auto-matched 61 of 71 devices using HR data. 10 remain unresolved.

A status bar at the top shows a progress indicator: "61 of 71 assigned" with an "Auto-match remaining" action button.

The main area shows an assignment table with columns: checkbox, device serial, assigned user (name and email), security profile, confidence level, and an action column (Accept / Change). Sample rows:

- 5CG4ABCD01 | Sara Mensah · sara.mensah@acme.com · Finance | FIN-SEC-STRICT | High | Accept · Change
- 5CG4ABCD05 | Liam O'Neill · Strategy | EXEC-PREMIUM | High | Accept · Change
- 5CG4ABCD10 | James Okafor · Engineering | DEV-ELEVATED | High | Accept · Change
- 5CG4ABCD61 | – | GENERAL-STD | Unresolved | Assign button
- 5CG4ABCD64 | – | GENERAL-STD | Unresolved | Assign button

Unresolved rows are visually distinct with a warning indicator.

A side panel shows "10 Devices Unresolved" with details:
5CG4ABCD61 (No HR match – serial not found in user export),
5CG4ABCD64 (Duplicate serial – in both Marketing and IT rosters),
5CG4ABCD72 (Missing owner attribute in CSV). Bulk resolve options:

assign to a department, defer resolution, or upload an updated HR roster.

The "Next: Review & Confirm" button advances to the final step.

SCREEN 6 of 8: Review & Confirm Deployment

Adam is on the final wizard step. All 4 progress steps are complete except step 4 (Review) which is active.

This is the final confirmation before deploying. An exception banner at the top warns: "4 items flagged – resolve before submitting."

A summary table shows all 71 devices with their assignment status. Flagged items (error serials 5CG4QRST55, 5CG4MNOP33) appear at the top with warning indicators. Clean assignments (Sara Mensah → 5CG4ABCD01 → FIN-SEC-STRICT, Liam O'Neill → 5CG4ABCD05 → EXEC-PREMIUM, etc.) follow.

A configuration panel on the right offers: Deploy immediately toggle (on by default), send notification emails to all 71 users (checkbox), and an estimated completion time of 4-6 hours.

"Submit Deployment" is the primary action. Cancel returns to the dashboard. If any flagged items are unresolved, the submit button shows a confirmation dialog: "4 devices have issues. Deploy anyway and handle these separately?"

SCREEN 7 of 8: Deployment Status Dashboard

The wizard is complete. Adam is now monitoring the deployment in real-time. This is a monitoring screen, not a wizard step – no progress indicator.

A breadcrumb shows: Fleet / Deployments / DEPL-2024-0847.

A prominent progress card at the top shows: 94% complete with a progress bar, 67 devices enrolled, 0 pending, 4 failed. Estimated completion: 14:30. An "Active Deployment" indicator and a timestamp ("Last updated 14:05") with a pause option.

Below, a filterable device table with tabs: All (71), Enrolled (67), Pending (0), Failed (4), Deferred (6). Default view shows failed devices first:

- 5CG4FAIL01 | james.okafor@acme.com | DEV-ELEVATED | Failed | 14:01 | "Profile conflict" | View details
- 5CG4FAIL02 | sara.lee@acme.com | SALES-WEST | Failed | 13:58 | "Incompatible OS" | View details
- 5CG4FAIL03 | mark.t@acme.com | CORP-GENERAL | Failed | 13:50 | "User declined" | View details
- 5CG4FAIL04 | admin.system@acme.com | INFRA-CORE | Failed | 13:40 | "Network unreachable" | Retry + View details

Then enrolled devices with their success timestamps.

Bottom bar: "Retry all failed (4)" secondary action, "Export status report", and "Complete deployment" primary (disabled while failures exist, tooltip: "4 devices still active").

SCREEN 8 of 8: Deployment Complete

All issues are resolved and Adam has marked the deployment complete. This is a confirmation and summary screen.

A centered success state: large checkmark, "Deployment complete" heading, with "DEPL-2024-0847 · Completed 14-Mar-2026 16:42 by Adam Chen" below it.

Three metric cards: 67 Enrolled, 4 Failed (resolved), 6 Deferred.

Details: "Asset inventory updated automatically" with a checkmark, and a link to "Download deployment report (PDF)."

Two actions: "View assignment records" (primary) and "Return to Fleet Dashboard" (secondary).

CONNECTION GUIDE

Screen	Action	Goes To
1. Fleet Dashboard	' + Add Devices'	2. Import
2. Import	'Next: Get AI Profile Recommendations'	3. AI Profiles
2. Import	'Cancel'	1. Fleet Dashboard
3. AI Profiles	'Next'	4. Review Profiles
3. AI Profiles	'Back'	2. Import
4. Review Profiles	'Next: Assign Users'	5. Assign Users
4. Review Profiles	'Back'	3. AI Profiles
5. Assign Users	'Next: Review & Confirm'	6. Review & Confirm
5. Assign Users	'Back'	4. Review Profiles
6. Review & Confirm	'Submit Deployment'	7. Deployment Status
6. Review & Confirm	'Cancel'	1. Fleet Dashboard
7. Deployment Status	'Complete deployment'	8. Deployment Complete
8. Complete	'View assignment records'	(Assignment Records)
8. Complete	'Return to Fleet Dashboard'	1. Fleet Dashboard

Example 2: Call Centre SaaS — Guided Topic Recommendation (8 screens, Tool-agnostic, Web)

This example uses the actual User Workflows output provided as input.

PRODUCT BRIEF

CallFlow – call centre topic management platform for support teams.
Platform: Web.

User: Casey (Call Centre Agent) – handles high-volume calls, needs to find the right topic quickly during live conversations.

Accuracy matters but speed matters more.

Style: Clean, efficient SaaS interface similar to Zendesk or Intercom. Minimal friction, fast-scanning layouts.

Design system is pre-loaded.

Navigation: Left sidebar with Dashboard, Topics, Calls, Analytics, Knowledge Base, Settings. Top bar with search and user avatar.
Topics is active for this flow.

Flow: Guided Topic Recommendation – 8 screens total (5 core + 1 popup + 2 error states). Agent finds the right call topic via

system-guided recommendations.

SCREEN 1 of 8: Dashboard

Casey logs into CallFlow and sees the main dashboard, which is her home screen for all administrative functions.

The sidebar shows Dashboard as the active section. The main area shows today's call statistics: 47 calls handled, 12 in queue, average handle time 4:32, and first-call resolution at 78%.

A "Quick Actions" section provides shortcuts to the most common tasks: Start New Call, Browse Topics, View Recent Calls, and Check Knowledge Base.

Below that, a recent activity feed showing the last 5 call resolutions with caller name, topic used, and outcome.

Casey clicks "Browse Topics" to navigate to the topic list.

SCREEN 2 of 8: Topic List

Casey is on the topic list, which shows all available call topics she can assign to conversations.

The page header reads "Topics" with a total count (142 topics). Two prominent action buttons at the top: "Search topics" (a standard search/filter approach) and "Let the system find the right topic" (the guided path, highlighted as recommended).

Below, a filterable table of topics with columns: Topic Name, Category, Last Updated, Usage Count (last 30 days), and Status (Active/Archived). Sample rows:

- Billing Dispute | Billing | 2 days ago | 234 uses | Active
- Account Lockout | Security | 1 day ago | 189 uses | Active
- Refund Request | Billing | 5 days ago | 156 uses | Active
- Password Reset | Security | 3 days ago | 312 uses | Active

- Shipping Delay | Logistics | 1 day ago | 98 uses | Active

Filter dropdowns for Category and Status. A search bar for topic name.

Casey clicks "Let the system find the right topic" to try the guided approach.

VARIATION – Encourage use of guided search

Screen 2: Topic List

This variation addresses: How encourage Casey to use guided search for faster topic identification?

Instead of showing two equal action buttons at the top, this version makes the guided path the default experience. When Casey arrives on this screen, a brief contextual banner appears: "You're on a live call about a billing issue – we've found 3 likely topics. View recommendations?" with "Yes, show me" as the primary action and "No, I'll search manually" as a subtle text link.

The topic table is still visible below but slightly de-emphasized. This steers Casey toward the faster guided path without removing her manual option.

SCREEN 3 of 8: Guided Topic Finder

Casey enters the guided topic finder. The system has analyzed her current call context (caller history, recent interactions, conversation keywords) and proactively suggests filters.

The screen shows a "smart context" card at the top: "Based on this call, we suggest:" followed by pre-applied filter chips – "Billing" category, "Active" status, "High usage" sort. Casey can tap any chip to remove it, or add her own filters.

Below the smart filters, a live-updating result count: "Showing 12 of 142 topics matching your context."

The filter controls are still visible and interactive – Casey can switch from the suggested "Billing" category to "Security" if the AI's context analysis was wrong. A "Reset to AI suggestions" link lets her go back to the system recommendations.

At the bottom: "Generate recommendations" button that runs the final matching algorithm and takes Casey to the results.

VARIATION – Leverage system intelligence to reduce manual filtering
Screen 3: Guided Topic Finder

This variation addresses: How leverage system intelligence to reduce manual filtering entirely?

Instead of showing filter controls with pre-applied suggestions, this version removes visible filters entirely. The system shows a natural language summary: "We analyzed this call and detected: billing concern, returning customer, previous dispute 30 days ago. Searching for the best matching topics..." with a brief loading animation, then immediately transitions to the recommended topics screen.

Casey sees a "Not quite right?" link that expands the manual filter controls if the auto-detection missed the mark. This is the highest-trust version of the guided experience.

VARIATION – Auto-detect call context
Screen 3: Guided Topic Finder

This variation addresses: What if the system could auto-detect call context and pre-filter topics?

Instead of requiring Casey to review and confirm suggested filters, this version works like a real-time assistant. While Casey is on the call, a sidebar widget continuously shows "Listening for context..." and updates with detected keywords and categories as the conversation progresses. When confidence exceeds 80%, it proactively highlights: "Billing Dispute topic matches this call (92% confidence)." Casey can confirm with one click without ever visiting a separate filter screen.

SCREEN 4 of 8: Recommended Topics

The system has processed the context and presents its top recommendations. Casey needs to pick the right topic for this call.

The screen shows a ranked list of 5 recommended topics, each as a card with:

- Topic name and category badge
- Match score (e.g., "92% match" / "85% match" / "71% match")
- A brief explanation of why this topic was recommended ("Caller has previous billing dispute, same issue category")
- Last used timestamp and usage frequency
- A "Select" button

The top recommendation (Billing Dispute, 92% match) is visually highlighted. Below the ranked list, a "None of these match – search all topics" link returns Casey to the full topic list.

Casey clicks "Select" on Billing Dispute to continue.

VARIATION – Make recommendations transparent and adjustable

Screen 4: Recommended Topics

This variation addresses: How ensure recommendations are transparent and adjustable?

Instead of showing recommendations as a simple ranked list, this version shows each recommendation with an expandable "Why this topic?" section. When expanded, it shows the data signals: "92% match based on: caller's last 3 calls were billing-related (40% weight), current conversation contains keywords 'charge', 'dispute', 'refund' (35% weight), this topic resolves similar calls in 3.2 minutes average (25% weight)."

Each recommendation also has a "Not relevant" button that adjusts the ranking in real-time and improves future recommendations.

SCREEN 5 of 8: Topic Details

Casey has selected "Billing Dispute" and is now reviewing its details before confirming.

The screen shows comprehensive topic information:

- Topic name: Billing Dispute
- Category: Billing
- Description: "For calls where the customer disputes a charge, requests investigation into a transaction, or escalates a billing concern."
- Resolution scripts: 3 scripts available (Quick resolution, Escalation path, Refund authorization)
- Average handle time: 6:45
- First-call resolution rate: 82%
- Related topics: Refund Request, Payment Failure, Charge Investigation

A prominent "Confirm Selection" button opens the final confirmation popup. A "Back to recommendations" link returns to the previous screen.

SCREEN 6 of 8: No Recommendations (Error State)

The system couldn't generate any recommendations for the current call context. This might happen when the call topic is unusual or the metadata is insufficient.

The screen shows a centered empty state with an illustration and the message: "We couldn't find matching topics for this call." Below: "This sometimes happens with new call types or when call context isn't clear enough."

Two recovery actions: "Search all topics manually" (primary) and "Retry recommendations" (secondary, tries again with relaxed matching criteria). A help link offers: "Need help? Contact your

team lead."

VARIATION – Reduce frustration when recommendations fail

Screen 6: No Recommendations

This variation addresses: How avoid user frustration when recommendations fail?

Instead of a binary "no results found" message, this version shows partial results: "We couldn't find a strong match, but here are 3 topics that partially fit:" followed by lower-confidence recommendations (e.g., 45%, 38%, 32% match scores). This gives Casey something to work with rather than a dead end. The "Search manually" option is still available but feels less like the only escape route.

SCREEN 7 of 8: Database Error (Error State)

The system experienced a database or connectivity error while loading topics or generating recommendations.

A full-screen error state with an apologetic message: "Something went wrong – we couldn't reach the topic database." Technical details are hidden behind a "Show details" expandable (error code, timestamp).

Two actions: "Try again" (primary, retries the last action) and "Contact support" (secondary, opens a support channel with the error details pre-filled). A note: "Your call data has been saved – you won't lose progress."

SCREEN 8 of 8: Confirmation Pop-Up (over Screen 5)

This appears as a modal dialog over the Topic Details screen when Casey clicks "Confirm Selection."

The modal shows: "Confirm topic selection?" with a summary: "You're assigning Billing Dispute to this call. This will:" followed by a checklist: load the resolution scripts, update the call record, and notify the quality team.

Two buttons: "Confirm" (primary) and "Cancel" (returns to Topic Details). A checkbox option: "Don't ask me again for this session."

CONNECTION GUIDE

Screen	Action	Goes To
-----	-----	-----
1. Dashboard	'Browse Topics'	→ 2. Topic List
2. Topic List	'Let the system find the right topic'	→ 3. Guided Finder
2. Topic List	'Search topics'	→ (manual filter on same screen)
3. Guided Finder	'Generate recommendations'	→ 4. Recommended Topics
3. Guided Finder	(if system fails)	→ 6. No Recommendations
3. Guided Finder	(if database error)	→ 7. Database Error
4. Recommended Topics	'Select' (any topic)	→ 5. Topic Details
4. Recommended Topics	'None of these match'	→ 2. Topic List
5. Topic Details	'Confirm Selection'	→ 8. Confirmation Pop-Up
5. Topic Details	'Back to recommendations'	→ 4. Recommended Topics
6. No Recommendations	'Search all topics manually'	→ 2. Topic List
6. No Recommendations	'Retry recommendations'	→ 3. Guided Finder
7. Database Error	'Try again'	→ (retry last action)
8. Confirmation Pop-Up	'Confirm'	→ 1. Dashboard (call updated)
8. Confirmation Pop-Up	'Cancel'	→ 5. Topic Details

Example 3: Healthcare — Patient Check-in (5 screens, Tool-agnostic, Mobile)

PRODUCT BRIEF

MedConnect – patient self check-in and health management app.
Platform: Mobile.
User: Priya Sharma (Patient) – busy professional, dislikes paperwork and phone-based admin. Wants the fastest path through

check-in.

Style: Warm, approachable healthcare app similar to Calm meets One Medical. Soft, reassuring, not clinical.

Design system is pre-loaded.

Navigation: Bottom tabs – Home, Appointments, Records, Profile.

Minimal top bar with app logo and notification bell.

Flow: Walk-in check-in – 5 screens. Confirm identity, verify insurance, optionally report symptoms, get checked in.

SCREEN 1 of 5: Home

Priya opens MedConnect and sees her home screen. The Home tab is active in the bottom navigation.

A prominent card shows today's appointment: "Dr. Anil Kapoor – Cardiology" at 10:15 AM, City Heart Hospital, Room 308. Two buttons: "Check in now" (primary, full-width) and "View details" (secondary).

Below, a 2x2 grid of quick actions: Book appointment, My records, Prescriptions, and Messages.

A "Recent" section shows: "Dr. Kapoor visit • 15-Jan-2026 • View summary →" and "Lab results • 08-Jan-2026 • View →".

Priya taps "Check in now" to begin.

SCREEN 2 of 5: Identity Confirmation

Priya has tapped "Check in now" and the app is confirming her identity. The top bar changes to show "← Check In" as a back link and "Step 1 of 3" on the right. A progress bar shows 33% complete. The Appointments tab is active.

The screen is simple and centered: a circular avatar with her initials "PS", her name "Priya Sharma", and her details: DOB 22 Sep 1991, ID PAT-10234, City Heart Hospital.

Below: "Is this you?" with "Please confirm your identity to continue."

Two full-width buttons at the bottom: "Yes, that's me" (primary) and "This isn't me" (secondary, returns to home).

SCREEN 3 of 5: Insurance Verification

Priya confirmed her identity and is now on step 2 of 3. Progress bar shows 66%.

Her insurance information is displayed in a card: HealthGuard Gold, Policy HG-44291, Status: Active, with a "Verified automatically" note.

Below the card: "Insurance details incorrect?" as a text link. Tapping it expands an inline form with two fields (Provider name and Policy number) and Update/Cancel buttons. This is collapsed by default – most patients won't need it.

A "Continue" button at the bottom advances to the symptom screen.

SCREEN 4 of 5: Symptom Input

Step 3 of 3. Progress bar shows 100% (this is the final step before confirmation).

A friendly heading: "Help Dr. Kapoor prepare for your visit."
A caption: "Optional – takes about 2 minutes."

A short form:

- Main concern (text area, "What brings you in today?")

- Severity (three options: Mild, Moderate, Severe – none pre-selected)
- Duration (dropdown: Today, 2-3 days, This week, Longer than a week)
- Any medications taken? (Yes / No radio). If Yes is selected, a text field appears: "e.g., Paracetamol 500mg"

Two actions: "Submit and check in" (primary) and "Skip and check in" (text link below – this entire form is optional).

SCREEN 5 of 5: Check-In Confirmed

Priya is checked in. The original top bar returns (app logo + notification bell). Appointments tab is active.

A centered success state: large checkmark icon, "You're checked in" heading, "Dr. Anil Kapoor · 10:15 AM today" below.

A wait info card: "8 minutes" in large text, "1 patient ahead of you", "We'll notify you when it's your turn."

A location card: pin icon, "Room 308 · Level 3", "City Heart Hospital", with a "Get directions" link.

Two buttons: "View appointment details" (secondary) and "Return to home" (text link).

CONNECTION GUIDE

Screen	Action	Goes To
1. Home	'Check in now'	→ 2. Identity
2. Identity	'Yes, that's me'	→ 3. Insurance
2. Identity	'This isn't me'	→ 1. Home
3. Insurance	'Continue'	→ 4. Symptoms
4. Symptoms	'Submit and check in'	→ 5. Confirmed
4. Symptoms	'Skip and check in'	→ 5. Confirmed

| 5. Confirmed | 'Return to home' | → 1. Home |
